

WORK EXPERIENCE

- *Data Scientist at 9D Technologies (Feb 2026–Now):*

- Architected and implemented the **multi-agent orchestration layer** for a Google ADK-based search system, enabling query decomposition, parallel specialist-agent execution, result aggregation, and final-answer synthesis across Google, Tavily, and Brave Search.
- Took the multi-agent search system from internal prototype to an **AWS Marketplace-ready deployment**, handling containerization, dependencies, configuration, tool integrations, edge cases, and packaging; also packaged a separate real-time ASR pipeline with speaker diarization and identification for Marketplace deployment.
- Built a systematic evaluation pipeline for **50+ quantized language models (GGUF)** for on-device Android deployment, combining manual testing, DeepEval with custom datasets and metrics, and 20+ public benchmarks across six domains; recommended ~25 models (4–5 per domain) ultimately used in the app across device tiers.
- Optimized **system prompts, sampling parameters, and on-device safety guardrails** for domain-specific assistants, developing and training a TensorFlow safety/disclosure classifier and deploying it as a TFLite model across education, coding, research, mathematics, healthcare, and finance use cases.
- Developed a **domain-specific speech-to-speech evaluation framework** for a female-health AI assistant, extending Eval-Bench with custom scenarios and evaluation metrics for medical safety, red-flag escalation, unsupported claims, conciseness, opening/closing quality, relevant follow-ups, and turn-taking latency; benchmarked PersonaPlex against GPT Realtime models and informed the selection of GPT Realtime 2.1 Mini for initial deployment.
- Conducted **GPU-hosting vs API cost analysis** for the speech-to-speech system, balancing model quality, latency, expected demand, and infrastructure costs to inform deployment strategy and potential future migration to self-hosted inference.

- *Data Scientist at Afiniti (Sep 2024–Feb 2026):*

- Optimized real-time **agent-customer pairing models** serving 200,000+ customers per week, using Bayesian agent segmentation models in Stan and LightGBM/XGBoost caller segmentation models to improve call-center matching.
- Consistently achieved **1.5x–2x target gains** in client-defined pairing performance, increasing sales and customer retention relative to clients' existing agent-customer pairing methods.
- Led **migration from R to Python for production pipelines**, enabling more extensive hyperparameter tuning, validation, visualization, and adoption of advanced ML approaches while maintaining existing production functionality.
- Built **real-time production monitoring systems** with SQL-based metrics, dashboards, and alerts to track pairing quality, agent utilization, caller wait times, and other operational indicators and proactively identify model or deployment issues.
- Led **ONNX pipeline development** for caller segmentation models, adapting model evaluation and production workflows to support ONNX-based inference for improved scalability.
- Collaborated with **cross-functional teams** to translate statistical and machine learning models into production call-center optimization systems and address real-world operational constraints.

- *Teaching Assistant at Lahore University of Management Sciences (Jan 2024–Jun 2024):* Co-instructed a **graduate-level course on Convex Optimization** with Dr. Hassan Mohy Ud Din.

RESEARCH EXPERIENCE

- *Under Dr. Kim-Chuan Toh (Oct 2025–Now):*

- Reimplemented the **ADMM component of the SDPNAL+ solver** from MATLAB to Julia for cross-language reproducibility and extensibility.
- Ported the solver from CPU to GPU using CUDA with **custom kernels and restructured data layouts** for efficient GPU utilization.
- Benchmarked implementations across 30+ test problems, **demonstrating 2–30x speedups** on GPU over CPU for medium-to-large instances.
- Restructured MATLAB-specific constructs (e.g., cell arrays, MEX dependencies) into Julia-native designs, improving **computational efficiency and maintainability**.
- Evaluated an alternative **PSD projection method** to replace the costly eigendecomposition-based approach.
- Diagnosed and resolved numerical instability in failed factorizations, identifying **ill-conditioning from extreme spectral scaling** as the root cause.

- *Under Dr. Hassan Mohy Ud Din (Jun 2024-Now):*
 - Conducted **end-to-end research**, including problem formulation, literature review, experimental design, and scientific writing ([Wave-GMS](#)).
 - Designed and implemented **computer vision models for medical image segmentation** using diffusion-based and variational approaches.
 - Built **training and evaluation pipelines in PyTorch** for systematic comparisons against strong baselines.
 - Analyzed results for **model architectures and loss functions** to improve segmentation performance.
- *Under Dr. Muhammad Tahir (Jun 2023-Jun 2024):*
 - Refined an **unrolled optimization algorithm for matrix completion** to improve baseline model ([convmc-net](#)).
 - Developed a novel extension to handle settings involving **Gaussian mixture model noise** ([ConvHuberMC-Net](#)).
 - Reproduced and evaluated prior methods, conducting **state-of-the-art benchmarking comparisons**.
 - Documented methodology and results in a **comprehensive technical report** ([report](#)).

PUBLICATIONS

- Talha Ahmed, **Nehal Ahmed Shaikh**, Hassan Mohy Ud Din. *Wave-GMS: Lightweight Multi-Scale Generative Model for Medical Image Segmentation*. ICASSP 2026.

PROJECTS

- **Speech Recognition and Translation System For Medical Communication:** Combined speech recognition, machine translation, and TTS functionality, selecting relevant vector databases, fine-tuning large language models (LLMs), and employing RAG pipelines to enable users to receive real-time state-of-the-art medical advice and information (see: [repository](#)).
- **Panel Data and Tobit Analysis on Health Care Dataset:** Analyzed healthcare service utilization in Germany employing tobit, fixed effects, and random effects models on a unique panel data set spanning from 1984 to 1995 (see: [data](#)), while focusing on socio-economic disparities in healthcare access and hence contributing critical insights to health-care economics. (see: [PDF](#)).
- **Air Pollution & Academic Performance:** Explored the relationship between air pollution and academic performance in Pakistan to know whether there exists any relationship between the two and to predict academic performance scores based on a customized synthesis of ASER (education) and NASA (climate) data (see: [blog](#)).
- **Sentiment Analysis on Audio Recordings:** Applied various classifiers (K-NN, logistic regression, naive Bayes, SVM, and neural network) on [CREMA-D](#) to predict the emotion signified by an audio and performed comparative metric analysis (see: [PDF](#)).
- **Regression Analysis to Estimate the Effect of Gender on Academic Performance:** Utilized a multiple linear regression model, featuring a host of explanatory variables, to conclude whether there is any significant difference between the CGPA values of male and female students, using field data collected through questionnaires (see: [PDF](#)).
- **Regression Analysis on the Determinants of New York Housing Prices:** Created a multiple linear regression model to determine the significance of factors affecting housing prices in New York and the effect of outliers by performing extensive data analysis (see: [data](#) and [PDF](#)).
- **Deep Learning for Computer Vision:** Implemented four deep learning algorithms—CNN, auto-encoder, LSTM, and ViT—for computer vision tasks (see: [repository](#)).

EDUCATION	BSc in Economics-Mathematics with a minor in Computer Science from Lahore University of Management Sciences.
COURSEWORK	Reinforcement Learning, Generative AI, Deep Learning, Machine Learning, Data Science, Signal Processing, Convex Optimization, Numerical Analysis, Differential Equations, Real Analysis, Statistics, Probability, Advanced Econometrics, Game Theory, Moral Philosophy.
HONORS	Placed four times on Dean's Honor List for excellent academic performance and ultimately graduated with distinction (3.72/4 CGPA).
SKILLS	Programming languages (Python, C/C++, R, SQL, Julia), machine learning (PyTorch, TensorFlow, LightGBM, XGBoost, Stan, ONNX), LLMs (Google ADK, DeepEval, llama.cpp, GGUF, TFLite, prompt engineering, LLM evaluation), tools (Docker, Git/GitHub, Linux, Superset, LaTeX).